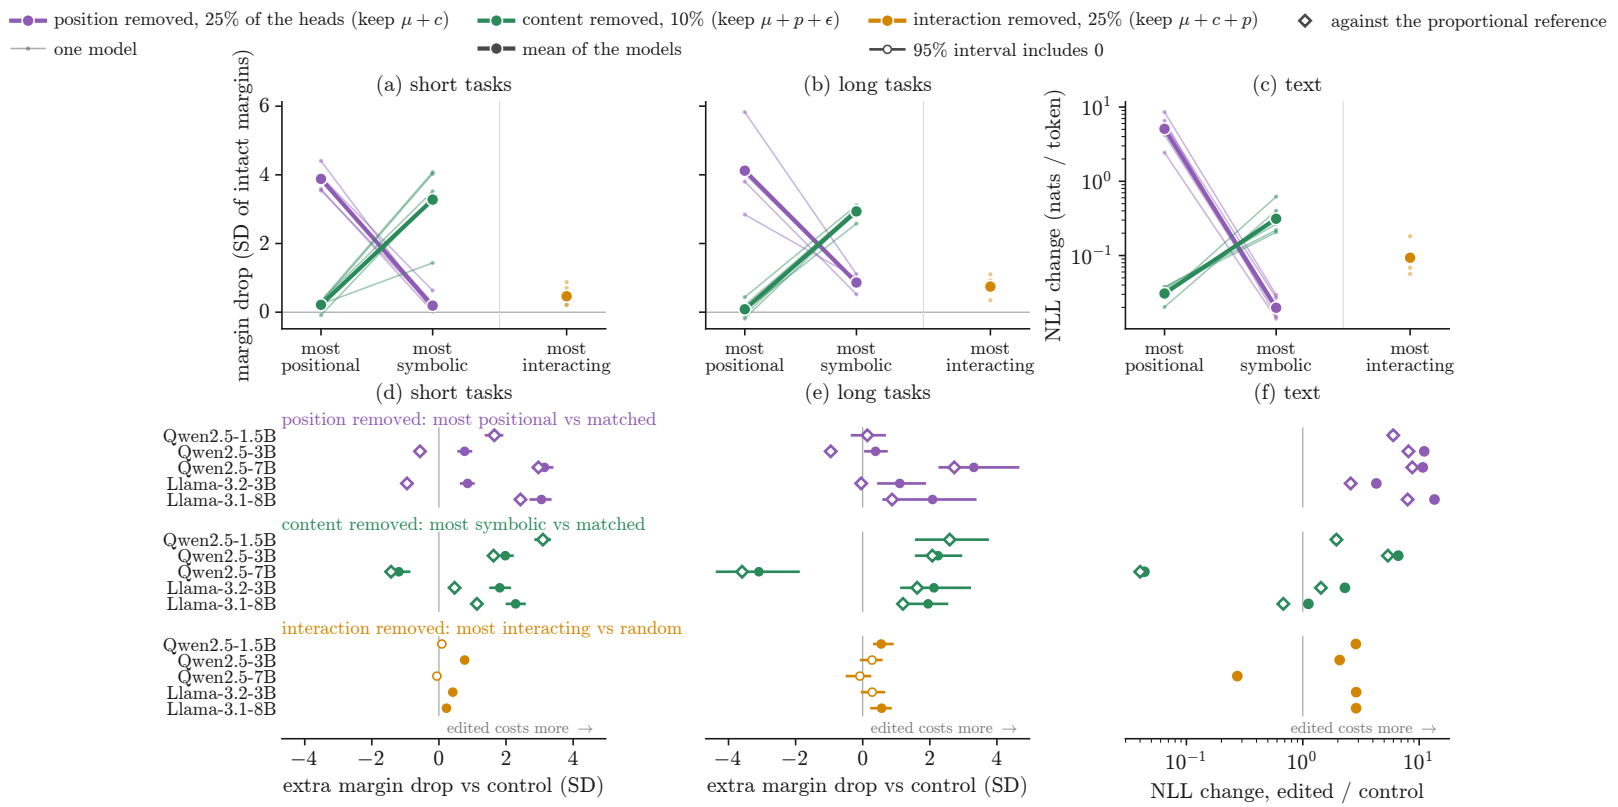
position removed, 25% of the heads (keep $\mu + c$)
content removed, 10% (keep $\mu + p + \epsilon$)
interaction removed, 25% (keep $\mu + c + p$)
against the proportional reference
one model
mean of the models
95% interval includes 0
(a) short tasks
(b) long tasks
(c) text
margin drop (SD of intact margins)
NLL change (nats / token)
most positional
most symbolic
most interacting
(d) short tasks
(e) long tasks
(f) text
position removed: most positional vs matched
content removed: most symbolic vs matched
interaction removed: most interacting vs random
Qwen2.5-1.5B
Qwen2.5-3B
Qwen2.5-7B
Llama-3.2-3B
Llama-3.1-8B
edited costs more →
extra margin drop vs control (SD)
NLL change, edited / control